

Networked systems – Sockets

Pr Sara Bouchenak
Sara.Bouchenak@insa-lyon.fr
<http://liris.cnrs.fr/~sbouchen/>



Examples of network-based applications

- Email services
 - e.g. Gmail
- Document sharing services
 - e.g. Dropbox
- Social media
 - e.g. Facebook, Twitter
- Video-on-demand services



OSI layers

Layer	Function	Examples of protocols
7. Application	Serves the actual applications used by the end user	SMTP, HTTP
6. Presentation	Formats data to be presented to the application layer	JPEG, GIF, ASCII
5. Session	Session establishment between processes running on different machines	RPC
4. Transport	Ensures that messages are delivered error-free, in sequence, and with no losses or duplication	TCP, UDP
3. Network	Controls the operations of the subnet, deciding which physical path the data takes	IP
2. Data link	Error-free transfer of data frames from one node to another over the Physical layer	PPP
1. Physical	Transmission and reception of unstructured raw bit stream over the physical medium	

OSI layers (cont.)

Layer	Function	Examples of protocols
7. Application	Serves the actual applications used by the end user	SMTP, HTTP
6. Presentation	Formats data to be presented to the application layer	JPEG, GIF, ASCII
5. Session	Session establishment between processes running on different machines	RPC
4. Transport	Ensures that messages are delivered error-free, in sequence, and with no losses or duplication	TCP, UDP
3. Network	Controls the operations of the subnet, deciding which physical path the data takes	IP
2. Data link	Error-free transfer of data frames from one node to another over the Physical layer	PPP
1. Physical	Transmission and reception of unstructured raw bit stream over the physical medium	

Objectives

- Introduce paradigms and mechanisms to build networked-based applications
 - Sockets
 - Remote Procedure Call (RPC / Java RMI)
 - World Wide Web
- Practical work
 - Application implementations

Layer	Function	Examples of protocols
7. Application	Services the actual applications used by the end user	SMTP, HTTP
6. Presentation	Formats data to be presented to the application layer	JPEG, GIF, ASCII
5. Session	Session establishment between processes running on different machines	RPC
4. Transport	Ensures that messages are delivered error-free, in sequence, and with no losses or duplication	TCP, UDP
3. Network	Controls the operations of the network, deciding which physical path the data takes	IP
2. Data link	Error-free transfer of data frames from one node to another over the physical layer	PPP
1. Physical	Transmission and reception of unstructured raw bit stream over the physical medium	

Professors

- Sara Bouchenak
 - Professor INSA Lyon – LIRIS Laboratory
 - CM, TP
- Frédérique Biennier
 - Professor INSA Lyon – LIRIS Laboratory
 - CM

Agenda

Lectures	Lab
Introduction to Sockets	
Introduction to RPC/RMI	
Introduction to WWW	Building Socket-based applications (Part I)
Building multi-tier web applications	Building Socket-based applications (Part II)
Network infrastructures	Building HTTP servers (Part I)
Network infrastructures	Building HTTP servers (Part II)

Outline

1. Introduction to sockets
2. **Addressing**
3. Point-to-point communication with TCP sockets
4. Point-to-point communication with UDP sockets
5. Group communication with multicast
6. Locating network resources

Client/server applications

- The server provides some service
 - Examples
 - processing database queries
 - sending out current stock prices
- The client uses the service provided by the server
 - Examples
 - displaying database query results to the user
 - making stock purchase recommendations to an investor


© S. Bouchenak

Networked systems

9

Client/server applications (2)

- Synchronous communication




© S. Bouchenak

Networked systems

10

Protocol layers



© S. Bouchenak

Networked systems

11

TCP, UDP, IP

- IP (Internet Protocol)
 - Corresponds to the network layer of the OSI model
 - Addressing, routing and transport of data packets
- TCP (Transmission Control Protocol)
 - A reliable protocol
 - Guarantees the delivery of data packets (i.e. data can not be dropped)
 - Puts data packets together in the right order (i.e. data must arrive on the client side in the same order in which the server sent it)
- UDP (User Datagram Protocol)
 - Is not reliable
 - More efficient
 - Used with applications which put high demands on transmission power, but can accept reductions in transmission quality (e.g. transmission of image/sound in real time)

© S. Bouchenak

Networked systems

12

Ports

- A computer connected to a TCP/IP network is identified by its IP address (unique identifier of a computer)
- When a process needs to be addressed via the network, it is identified by
 - the underlying computer's IP address
 - + a port number associated with the process (unique identifier of a process)
- Port numbers are managed by the operating system
- Port numbers are allocated to processes when they want to use the network
- Many important services have a standardized port
 - Example: port 25 for telnet service


© S. Bouchenak

Networked systems

13

Ports (2)

- Server side
 - When a server process is started on a computer, the OS links it to a fixed port
 - The server waits as a background process (daemon) for incoming connections addressed to this port
- Client side
 - When a connection to a service is to be opened between a client process on computer and a server process on another computer, the client process is assigned a port number
 - The server process can then uniquely identify the client process in connection in the entire network
- Ports as communication points
 - Ports represent communication points



© S. Bouchenak

Networked systems

14

Sockets

- A socket may be seen as a communication channel between two points (i.e. two ports)
- Strictly speaking, a socket is an endpoint of a communication connection between two computers
- A socket is identified by an IP address and a port number
- From the programmer's point of view, a socket represents the mechanism to transfer data from one computer to another

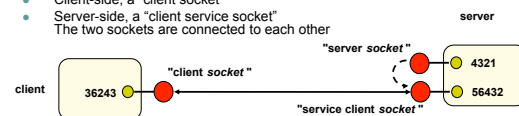
© S. Bouchenak

Networked systems

15

How sockets work

- Example with TCP/IP, in 3 steps:
 1. The server creates a "server socket" (associated with a port), and waits for incoming connections
 2. The client connects to the "server socket"; two sockets are created
 - Client-side, a "client socket"
 - Server-side, a "client service socket"
 - The two sockets are connected to each other
 3. The client and server communicate via the two sockets (writing/reading bytes); the "server socket" can accept new connections



© S. Bouchenak

Networked systems

16

Sockets – A bit of history

- Originally
 - Sockets were developed for BSD Unix, in the 1980s
 - Sockets used to be part of the operating system; they had to be invoked via system-specific libraries for C/C++
 - Programming distributed applications was hard (access was different from one OS to another, programs were not portable)
- Today
 - Sockets are available on all platforms and represent the most fundamental communication mechanism
 - Example: the Java programming interface for sockets abstracts them from the underlying OS, making them easier to use

Outline


1. Introduction to sockets
2. **Addressing**
3. Point-to-point communication with TCP sockets
4. Point-to-point communication with UDP sockets
5. Group communication with multicast
6. Locating network resources

Addressing in Java

- The *java.net* package provides the following addressing-related classes:
 - InetAddress
 - Inet4Address
 - Inet6Address
 - SocketAddress
 - InetSocketAddress

IP addressing in Java

- For IP addressing, three classes are provided:
 - **InetAddress** represents an IP address, which is either a 32- or 128-bit unsigned number used by IP
 - **Inet4Address** represents a 32-bit IPv4 address. It has the familiar form *n.n.n.n*, where *n* is a byte; e.g., 129.250.35.250
 - **Inet6Address** represents a 128-bit IPv6 address



Java classes related to IP addressing

- `java.net.InetAddress` class
 - Represents an IP address, which is either a 32- or 128-bit unsigned number used by IP
 - Textual representation of an IP address: it is family specific (see `Inet4Address` for example)
 - Host name associated with an IP address, e.g. *hoff.e.ejf-grenoble.fr*
 - Host name to IP address resolution: using a network naming service such as DNS (Domain Name System), NIS (Network Information Service)
 - Some API elements
 - `String getHostAddress()` Returns the IP address string in textual presentation.
 - `String getHostName()` Gets the host name for this IP address.
 - static `InetAddress getLocalHost()` Returns the local host.

© S. Bouchenak

Networked systems

21

Java classes related to IP addressing (2)

- `java.net.Inet4Address` class
 - Textual representation of IPv4 address has the familiar form *n.n.n.n*, where *n* is a byte; e.g., 129.250.35.250
 - Some API elements
 - `byte[] getAddress()` Returns the raw IP address of this `InetAddress` object.
 - `String getHostAddress()` Returns the IP address string in textual presentation form.


© S. Bouchenak

Networked systems

22

Socket addressing in Java

- Two classes are provided for socket addressing:
 - `SocketAddress`
 - It is an abstract socket address, independent of a specific protocol
 - It is intended for subclassing for a specific protocol (e.g. `InetSocketAddress`)
 - `InetSocketAddress`
 - It represents an IP socket address and can include:
 - an IP address (e.g., 129.250.35.250) and port (e.g., 80)
 - a hostname (e.g., coastnews.com) and port (e.g., 1000)
 - port only (e.g., 1010), in which case, a wildcard IP address is assumed



© S. Bouchenak

Networked systems

23

Outline

1. Introduction to sockets
2. Addressing
3. **Point-to-point communication with TCP sockets**
4. Point-to-point communication with UDP sockets
5. Group communication with multicast
6. Locating network resources

© S. Bouchenak

Networked systems

24

Java sockets over TCP

- Network communication in the connected mode (using TCP)
 - General schema:
 - a connection is opened between a client and a server,
 - a series of requests (i.e. messages) are exchanged between the client and the server,
 - the connection is closed
 - The server maintains a client session in which the state between different requests is maintained
 - TCP guarantees: reliability, delivery and order of messages
 - Adequate to long communication where several messages are exchanged (e.g. HTTP client sessions)

© S. Bouchenak

Networked systems

25

Java classes related to TCP sockets

- The *java.net* package provides the following classes related to TCP communication:
 - `ServerSocket`
 - `Socket`
- *ServerSocket*
 - It represents the socket on a server that waits and listens for requests for service from a client
- *Socket*
 - It represents the endpoints for communication between a server and a client.

© S. Bouchenak

Networked systems

26

Example of Java sockets/TCP

Server side	Client side
<pre>import java.net.*; int port = 1234; // Create a server socket associated with // the server port ServerSocket serverSocket = new ServerSocket(port); // End-less loop while (true) { System.out.println("Waiting for client_"); // Server waits for a connection Socket serverToClient = serverSocket.accept(); // A client connected System.out.println("Client " + client.getInetAddress() + " connected."); // Server receives a message from // client ... }</pre>	<pre>import java.net.*; String serverHost = "aun"; int serverPort = 1234; // Client connects to server Socket clientToServer = new Socket(serverHost, serverPort); // Client connected System.out.println("Connected to " + server.getInetAddress()); // Client sends a message to server ...</pre>

© S. Bouchenak

Networked systems

27

Streams


- Definition
 - Streams are an abstraction for arbitrary data streams
- Examples
 - Streams from/to a socket
 - Streams from/to a file
 - Streams from/to the console
- Input and output streams
 - Input streams used to receive (i.e. read) bytes
 - Output streams used to send (i.e. write) bytes

© S. Bouchenak

Networked systems

28

Client/server streams



© S. Bouchenak

Networked systems

29

Streams in Java

- The `java.io` package contains the following classes related to streams:
 - `InputStream` / `OutputStream`**
 - abstract classes that respectively represent input streams of bytes and output streams of bytes
 - `ObjectInputStream` / `ObjectOutputStream`**
 - classes respectively representing input streams of Java objects and output streams of Java objects
 - `FileInputStream` / `FileOutputStream`**
 - classes representing input streams for respectively reading data from a file or writing data to a file
 - `FilterInputStream` / `FilterOutputStream`**
 - classes that respectively contain other input streams or output streams, which it uses to possibly transform data along the way

© S. Bouchenak

Networked systems

30

Java sockets and streams

Server side	Client side
<pre>import java.io.*; ... // End-less loop while (true) { System.out.println("Waiting for client."); // Server waits for a connection Socket serverToClient = serverSocket.accept(); // A client connected System.out.println("Client " + serverToClient.getInetAddress() + " connected."); // Get the server's output stream OutputStream os = serverToClient.getOutputStream(); // Build data and transform it // to bytes Date date = new Date(); byte[] b = date.toString().getBytes(); // Write in output stream os.write(b); }</pre>	<pre>import java.io.*; ... // Client connects to server Socket clientToServer = new Socket(serverHost, serverPort); // Client connected System.out.println("Connected to " + clientToServer.getInetAddress()); // Get the client's input stream InputStream is = clientToServer.getInputStream(); // Read data from input stream byte[] b = new byte[100]; int num = is.read(b); // Transform data from bytes to String String date = new String(b); System.out.println("Server said: " + date);</pre>

© S. Bouchenak

Networked systems

31

Other Java stream types

Server side	Client side
<pre>import java.io.*; ... // End-less loop while (true) { System.out.println("Waiting for client."); // Server waits for a connection Socket serverToClient = serverSocket.accept(); // A client connected System.out.println("Client " + serverToClient.getInetAddress() + " connected."); // Get the server's output stream OutputStream os = serverToClient.getOutputStream(); // Get the associated object output // stream ObjectOutputStream oos = new ObjectOutputStream(os); // Write object in output stream Date date = new Date(); oos.writeObject(date); // Close output stream oos.close(); }</pre>	<pre>import java.io.*; ... // Client connects to server Socket clientToServer = new Socket(serverHost, serverPort); // Client connected System.out.println("Connected to " + clientToServer.getInetAddress()); // Get the client's input stream InputStream is = clientToServer.getInputStream(); // Get the associated object input stream ObjectInputStream ois = new ObjectInputStream(is); // Read object from input stream Date date = (Date) ois.readObject(); System.out.println("Server said: " + date); // Close input stream ois.close();</pre>

© S. Bouchenak

Networked systems

32

Outline

1. Introduction to sockets
2. Addressing
3. Point-to-point communication with TCP sockets
4. **Point-to-point communication with UDP sockets**
5. Group communication with multicast
6. Locating network resources

Java sockets over UDP

- Networking in the unconnected mode (using UDP)
 - Some applications that communicate over the network do not require reliable, point-to-point channel provided by TCP
- Applications might benefit from a mode of communication that delivers independent packages of information whose arrival and order of arrival are not guaranteed
- UDP protocol provides a mode of network communication whereby applications send packets of data, called datagrams, to one another.
- A datagram is an independent, self-contained message sent over the network whose arrival, arrival time, and content are not guaranteed.

Java classes related to UDP

- The *java.net* package provides the following classes related to UDP communication :
 - DatagramPacket
 - DatagramSocket
- **DatagramPacket**
 - It represents a datagram packet used for connectionless delivery and normally includes destination address and port information
- **DatagramSocket**
 - It is a socket used for sending and receiving datagram packets over a network via UDP
 - A *DatagramPacket* is sent from a *DatagramSocket* by calling the *send* method
 - The *receive* method is used for receiving a *DatagramPacket* on a *DatagramSocket*

Example of Java sockets/UDP

Server side

```
import java.net.*;

int port = 1234;
// Create a datagram socket associated
// with the server port
DatagramSocket serverSock = new
    DatagramSocket(port);
// End-less loop
while (true) {
    System.out.println("Waiting for client
    packet...");
    byte[] buf = new byte[256];
    // Create a datagram packet
    DatagramPacket packet = new
        DatagramPacket(buf, buf.length);
    // Wait for a packet
    serverSock.receive(packet);
    // Get client IP address and port number
    InetAddress clientAddr = packet.getAddress();
    int clientPort = packet.getPort();
    // Build a response
    initialize buf ...
    // Build a datagram packet for response
    packet = new DatagramPacket(buf, buf.length,
        clientAddr, clientPort);
    // Send a response
    serverSock.send(packet);
}
```

Client side

```
import java.net.*;

int serverPort = 1234;
String serverHost = "...";
// Create a datagram socket
DatagramSocket clientSock = new DatagramSocket();
byte[] buf = new byte[256];
// Get server's IP address
InetAddress serverAddr =
    InetAddress.getByName(serverHost);
// Build a request
initialize buf ...
// Create a datagram packet destined for the
// server
DatagramPacket packet = new DatagramPacket(buf,
    buf.length, serverAddr, serverPort);
// Send datagram packet to server
clientSock.send(packet);
// Build a datagram packet for response
packet = new DatagramPacket(buf, buf.length);
// Receive response
clientSock.receive(packet);
String received = new String(packet.getData(), 0,
    packet.getLength());
System.out.println("Response: " + received);
```

Outline

1. Introduction to sockets
2. Addressing
3. Point-to-point communication with TCP sockets
4. Point-to-point communication with UDP sockets
5. **Group communication with multicast**
6. Locating network resources

Multicast

- Besides point-to-point communication, Java provides multicast communication which can be used to communicate with a group of communication points
- Multicast is already available at the Internet protocol level
- Multicast is based on the UDP transport protocol
- Multicast communication uses IP addresses of a particular class (class D) which are not linked to host computers, but are exclusively reserved for multicast communication

Multicast – How it works

- First, all participants in a communication group must register with the group to be able to join the communication
- Messages can be sent to all group members or received from others
- To end the membership in a group, the participant must sign out
- Applications: multimedia conferences on the Internet are examples of applications for the multicast

Java multicast

- The *java.net* package provides the following classes related to multicast communication :
 - DatagramPacket
 - MulticastSocket
- **MulticastSocket**
 - A *MulticastSocket* is a (UDP) *DatagramSocket*, with additional capabilities for joining "groups" of other multicast hosts on the internet
 - One would join a multicast group by first creating a *MulticastSocket* with the desired port, then invoking the *joinGroup(InetAddress groupAddr)* method
 - When one sends a message to a multicast group, all subscribing recipients to that host and port receive the message
 - The socket does not need to be a member of the multicast group to send messages to it
 - When a socket subscribes to a multicast group/port, it receives datagrams sent by other hosts to the group/port, as do all other members of the group and port.
 - A socket relinquishes membership in a group by the *leaveGroup(InetAddress addr)* method

Java multicast example

```
import java.net.*;

...
// Group IP address
InetAddress groupAddr =
    InetAddress.getByAddress("228.5.6.7");
int groupPort = 6789;

// Create a multicast socket
MulticastSocket s = new
    MulticastSocket(groupPort);
// Join the group
s.joinGroup(groupAddr);

// Build a datagram packet for a message
// to send to the group
String msg = "Hello";
DatagramPacket hl = new
    DatagramPacket(msg.getBytes(),
        msg.length(), groupAddr, groupPort);

// Send a multicast message to the group
s.send(hl);

// Build a datagram packet for response
byte[] buf = new byte[1000];
DatagramPacket recv = new
    DatagramPacket(buf, buf.length);

// Receive a datagram packet response
s.receive(recv);
...
// OK, I'm done talking - leave the group
s.leaveGroup(groupAddr);
```

© S. Bouchenak

Networked systems

41

Outline

1. Introduction to sockets
2. Addressing
3. Point-to-point communication with TCP sockets
4. Point-to-point communication with UDP sockets
5. Group communication with multicast
6. Locating network resources

© S. Bouchenak

Networked systems

42

Identifying and locating network resources

- Examples of network resources
 - An image, audio file available on the network
 - A program file, such as a Servlet, available on the network
- Classes related to locating or identifying network resources (c.f. package *java.net*)
 - URI
 - URL
 - URLClassLoader
 - URLConnection
 - URLStreamHandler
 - HttpURLConnection
 - JarURLConnection


© S. Bouchenak

Networked systems

43

Identifying and locating network resources (2)

- **URI**
 - It represents a Uniform Resource Identifier for a resource.
 - It is an identifier for a resource but not necessarily a locator for that resource.
- **URL**
 - It represents a Uniform Resource Locator for a resource.
 - A *URL* tells how to access the resource, while a *URI* may or may not.
 - The protocol used to locate the resource is known from the *URL*.



© S. Bouchenak

Networked systems

44

URL – Uniform Resource Location



- A URL is a pointer to a "resource" on the World Wide Web
- A resource can be something as simple as a file or a directory, or it can be a reference to a more complicated object, such as a query to a database or to a search engine
- Example of a URL:
 - `http://liris.cnrs.fr`
- A URL has two main components:
 - Protocol identifier
 - Resource name

© S. Bouchenak

Networked systems

45

URL



- URL's protocol
 - The protocol identifier indicates the name of the protocol to be used to fetch the resource.
 - The example uses the Hypertext Transfer Protocol (HTTP), which is typically used to serve up hypertext documents.
 - HTTP is just one of many different protocols used to access different types of resources on the net.
 - Other protocols include File Transfer Protocol (FTP), File, and News.
- URL's resource
 - The format of the resource name depends entirely on the protocol used.
 - For many protocols, including HTTP, the resource name contains one or more of the components listed below:
 - **Host Name** The name of the machine on which the resource lives.
 - **Port Number** The port number to which to connect (typically optional).
 - **Filename** The pathname to the file on the machine.
 - **Reference** A reference to a named anchor within a resource that usually identifies a specific location within a file (typically optional).

© S. Bouchenak

Networked systems

46

Identifying and locating network resources in Java



- **URLConnection**
 - It is the abstract superclass of all classes that represent a connection between an application and a network resource identified by a URL.
 - Given a URL and hence a protocol, `URL.openConnection()` returns an instance of the appropriate implementation of `URLConnection` for the protocol.
 - The instance provides the means (`URLConnection.connect()`) to actually open the connection and access the URL.
- **HttpURLConnection**
 - It is the most commonly used implementation of `URLConnection`.
 - It is for http protocol, the protocol used for accessing content on web servers.

© S. Bouchenak

Networked systems

47

Outline



1. Introduction to sockets
2. Addressing
3. Point-to-point communication with TCP sockets
4. Point-to-point communication with UDP sockets
5. Group communication with multicast
6. Locating network resources

© S. Bouchenak

Networked systems

48

Agenda

Lectures	Lab
Introduction to Sockets	
Introduction to RPC/RMI	
Building multi-tier web applications	Building Socket-based applications (Part I)
	Building Socket-based applications (Part II)
Network infrastructures	Building HTTP servers (Part I)
Network infrastructures	Building HTTP servers (Part II)

References

This lecture is built from:

- Sun Microsystems. *Java Tutorial on Networking*. <http://java.sun.com/docs/books/tutorial/networking/>
- <http://www.reddit.com/r/networking/>
- M. Boger. *Java in Distributed Systems: Concurrency, Distribution and Persistence*. Wiley, 2001.
- This lecture is partly based on lectures given by Sacha Krakowiak, <http://sardes.inrialpes.fr/people/krakowia/>